



Boost UDP Transaction Performance

Toshiaki Makita
NTT Open Source Software Center

Today's topics



- **Background**
- **Basic technologies for network performance**
- **How to improve UDP performance**

Who is Toshiaki Makita?



- **Linux kernel engineer at NTT Open Source Software Center**
- **Technical support for NTT group companies**
- **Active patch submitter on kernel networking subsystem**

Background

UDP transactions in the Internet



- **Services using UDP**

- DNS
- RADIUS
- NTP
- SNMP
- ...

- **Heavily used by network service providers**

Ethernet Bandwidth and Transactions



- **Ethernet bandwidth evolution**

- 10M -> 100M -> 1G -> 10G -> 40G -> 100G -> ...
- 10G (or more) NICs are getting common on commodity servers

- **Transactions in 10G network**

- In the shortest packet case:
 - Maximum 14,880,952 packets/s^{*1}
- Getting hard to handle in a single server...

^{*1} shortest ethernet frame size 64bytes + preamble+IFG 20bytes = 84 bytes = 672 bits
 $10,000,000,000 / 672 = 14,880,952$

How many transactions to handle?

- **UDP payload sizes**

- DNS

- A/AAAA query: 40~ bytes
 - A/AAAA response: 100~ bytes

- RADIUS

- Access-Request: 70~ bytes
 - Access-Accept: 30~ bytes
 - Typically 100~ bytes with some attributes

- In many cases 100~ bytes

- **100 bytes transactions in 10G network**

- Max 7,530,120 transactions/s^{*1}
 - Less than shortest packet case, but still challenging

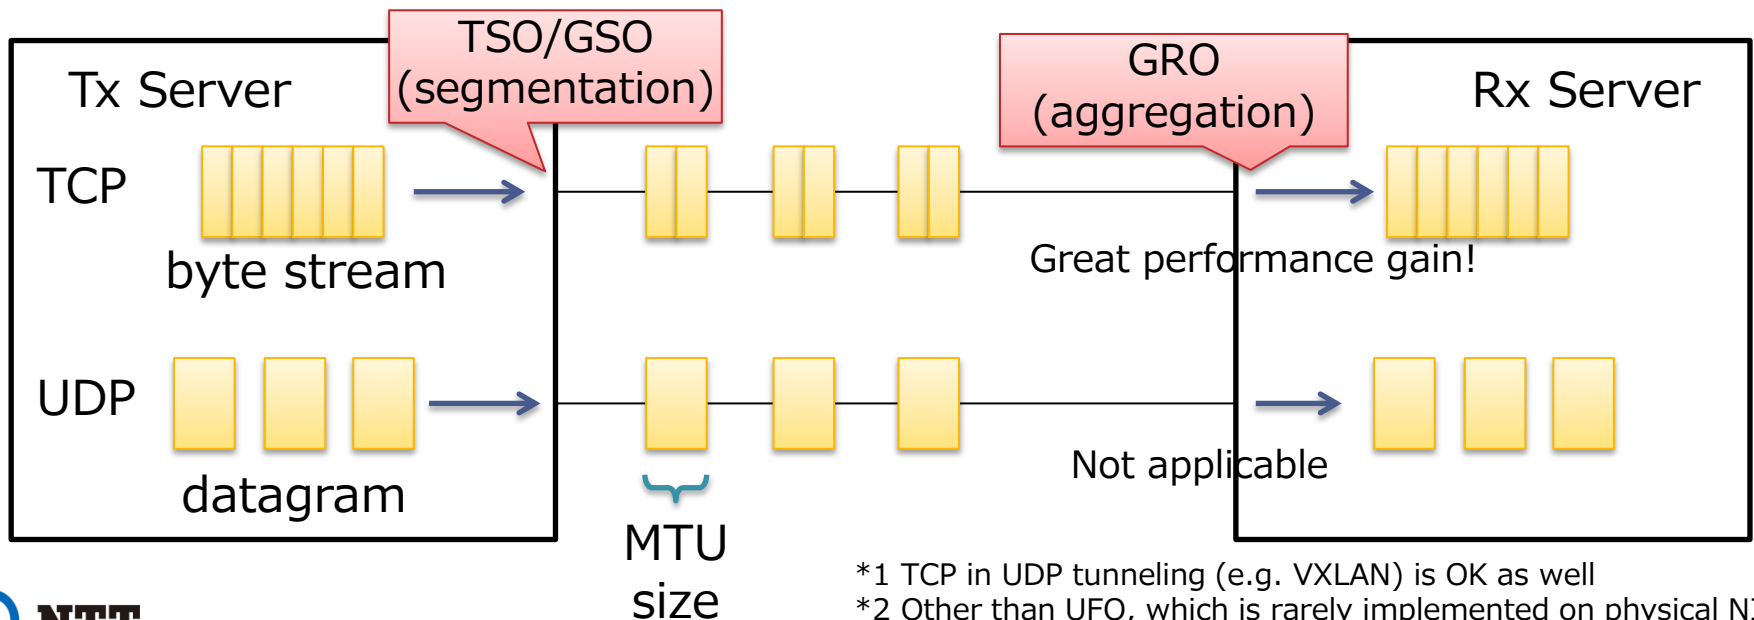
^{*1} 100 bytes + IP/UDP/Ether headers 46bytes + preamble+IFG 20bytes = 166 bytes = 1328 bits
10,000,000,000 / 1328 = 7,530,120

Basic technologies for network performance (not only for UDP)

Basic technologies for network performance

• TSO/GSO/GRO

- Packet segmentation/aggregation
- Reduce packets to process within server
- Applicable to TCP*¹ (byte stream)
- Not applicable to UDP*² (datagram) 😞
 - UDP has explicit boundary between datagrams
 - Cannot segment/aggregate packets

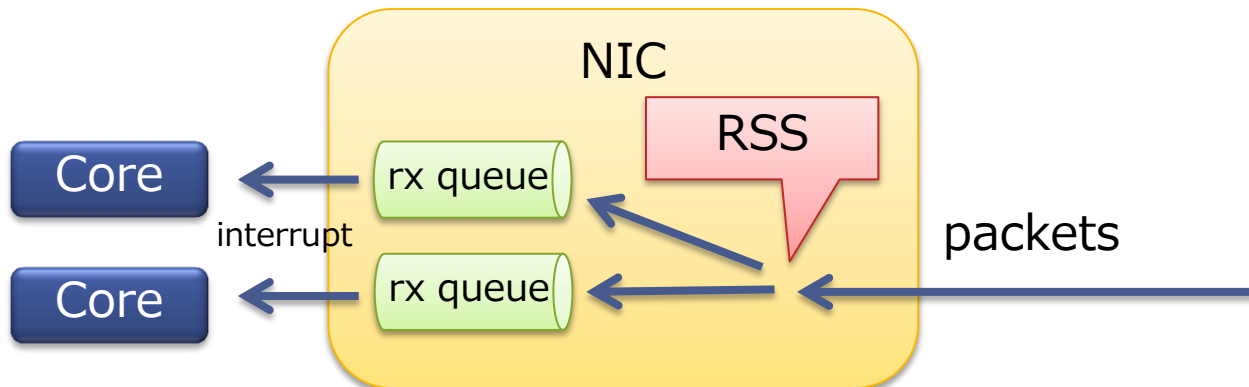


*1 TCP in UDP tunneling (e.g. VXLAN) is OK as well

*2 Other than UFO, which is rarely implemented on physical NICs

• RSS

- Scale network Rx processing in multi-core server
- RSS itself is a NIC feature
 - Distribute packets to multi-queue in a NIC
 - Each queue has a different interrupt vector
(Packets on each queue can be processed by different core)
- Applicable to TCP/UDP 😊
- Common 10G NICs have RSS




Performance with RSS enabled NIC



• 100 bytes UDP transaction performance

- Measured by simple^{*1} (multi-threaded) echo server
- OS: **kernel 4.6.3** (in RHEL 7.2 environment)
- Mid-range commodity server with **20 cores** and **10G NIC**:
 - NIC: Intel 82599ES (has RSS, **max 64 queues**)
 - CPU: Xeon E5-2650 v3 (2.3 GHz 10 cores) * 2 sockets
Hyper-threading off (make analysis easy, enabled later)
- Results: **270,000** transactions/s (tps) (approx. **360Mbps**)
 - **3.6%** utilization of 10G bandwidth



*1 create as many threads as core num
each thread just calls recvfrom() and sendto()

*2 There is only 1 client (IP address).

To spread UDP traffic on the NIC, RSS is configured
to see UDP port numbers. This setting is not
needed for common UDP servers.



NTT

Copyright © 2016 NTT Corp. All Rights Reserved. **11**

How to improve this?

Identify bottleneck

• sar -u ALL -P ALL 1

19:57:54	CPU	%usr	%nice	%sys	%iowait	%steal	%irq	%soft	%guest	%gnice	%idle
19:57:54	all	0.37	0.00	42.58	0.00	0.00	0.00	50.00	0.00	0.00	7.05
19:57:54	0	0.00	0.00	0.00	0.00	0.00	0.00	100.00	0.00	0.00	0.00
19:57:54	1	0.00	0.00	0.00	0.00	0.00	0.00	100.00	0.00	0.00	0.00
19:57:54	2	0.00	0.00	0.00	0.00	0.00	0.00	100.00	0.00	0.00	0.00
19:57:54	3	0.00	0.00	0.00	0.00	0.00	0.00	100.00	0.00	0.00	0.00
19:57:54	4	0.00	0.00	0.00	0.00	0.00	0.00	100.00	0.00	0.00	0.00
19:57:54	5	1.82	0.00	83.64	0.00	0.00	0.00	0.00	0.00	0.00	14.55
19:57:54	6	0.00	0.00	87.04	0.00	0.00	0.00	0.00	0.00	0.00	12.96
19:57:54	7	0.00	0.00	85.19	0.00	0.00	0.00	0.00	0.00	0.00	14.81
19:57:54	8	0.00	0.00	85.45	0.00	0.00	0.00	0.00	0.00	0.00	14.55
19:57:54	9	0.00	0.00	85.19	0.00	0.00	0.00	0.00	0.00	0.00	14.81
19:57:54	10	0.00	0.00	0.00	0.00	0.00	0.00	100.00	0.00	0.00	0.00
19:57:54	11	0.00	0.00	0.00	0.00	0.00	0.00	100.00	0.00	0.00	0.00
19:57:54	12	0.00	0.00	0.00	0.00	0.00	0.00	100.00	0.00	0.00	0.00
19:57:54	13	0.00	0.00	0.00	0.00	0.00	0.00	100.00	0.00	0.00	0.00
19:57:54	14	0.00	0.00	0.00	0.00	0.00	0.00	100.00	0.00	0.00	0.00
19:57:54	15	1.82	0.00	83.64	0.00	0.00	0.00	0.00	0.00	0.00	14.55
19:57:54	16	0.00	0.00	87.04	0.00	0.00	0.00	0.00	0.00	0.00	12.96
19:57:54	17	1.82	0.00	83.64	0.00	0.00	0.00	0.00	0.00	0.00	14.55
19:57:54	18	0.00	0.00	85.45	0.00	0.00	0.00	0.00	0.00	0.00	14.55
19:57:54	19	0.00	0.00	85.45	0.00	0.00	0.00	0.00	0.00	0.00	14.55

Node 0

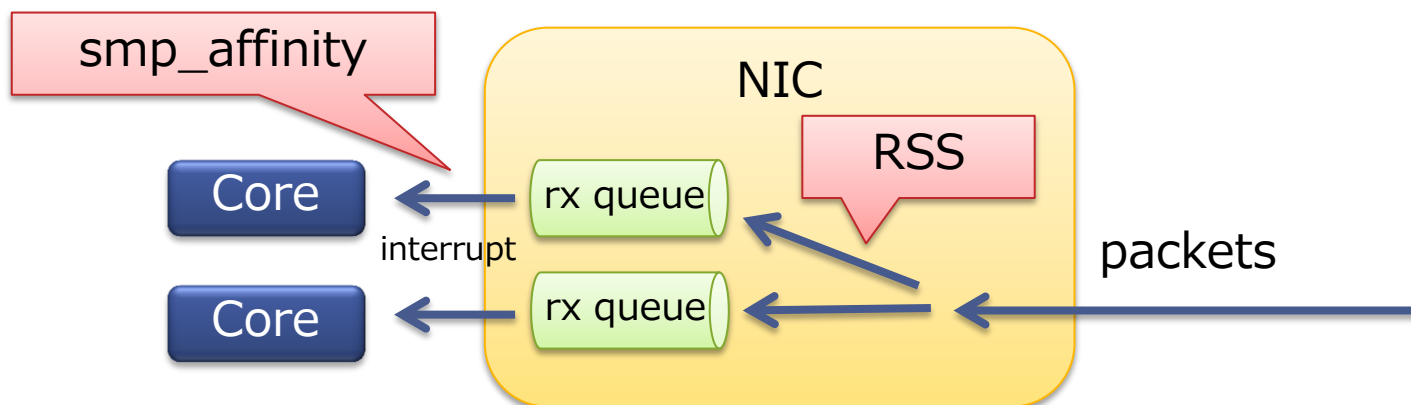
Node 1

• softirq (interrupt processing) is performed only on NUMA Node 0, why?

- although we have enough (64) queues for 20 cores...

softirq (interrupt processing) with RSS

- RSS distributes packets to rx-queues
- Interrupt destination of each queue is determined by `/proc/irq/<irq>/smp_affinity`




- `smp_affinity` is usually set by `irqbalance` daemon

Check smp_affinity

• smp_affinity*1

```
$ for ((irq=105; irq<=124; irq++)); do
>   cat /proc/irq/$irq/smp_affinity
> done
```

```
01000    -> 12    -> Node 0
00800    -> 11    -> Node 0
00400    -> 10    -> Node 0
00400    -> 10    -> Node 0
01000    -> 12    -> Node 0
04000    -> 14    -> Node 0
00400    -> 10    -> Node 0
00010    -> 4     -> Node 0
00004    -> 2     -> Node 0
02000    -> 13    -> Node 0
```



```
04000    -> 14    -> Node 0
00001    -> 0     -> Node 0
02000    -> 13    -> Node 0
01000    -> 12    -> Node 0
00008    -> 3     -> Node 0
00800    -> 11    -> Node 0
00800    -> 11    -> Node 0
04000    -> 14    -> Node 0
00800    -> 11    -> Node 0
02000    -> 13    -> Node 0
```

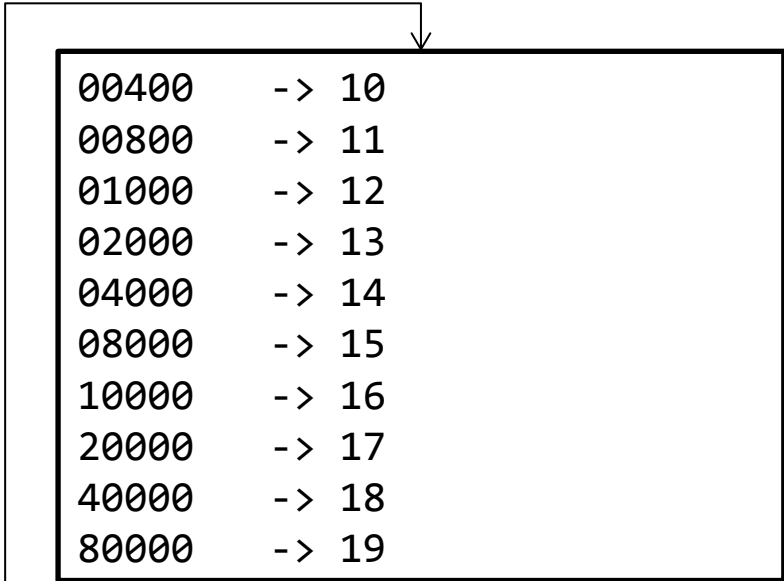
• irqbalance is using only Node 0 (cores 0-4, 10-14)

• Can we change this?

Check affinity_hint

- Some NIC drivers provide affinity_hint

```
$ for ((irq=105; irq<=124; irq++)); do  
>   cat /proc/irq/$irq/affinity_hint  
> done  
00001    -> 0  
00002    -> 1  
00004    -> 2  
00008    -> 3  
00010    -> 4  
00020    -> 5  
00040    -> 6  
00080    -> 7  
00100    -> 8  
00200    -> 9
```



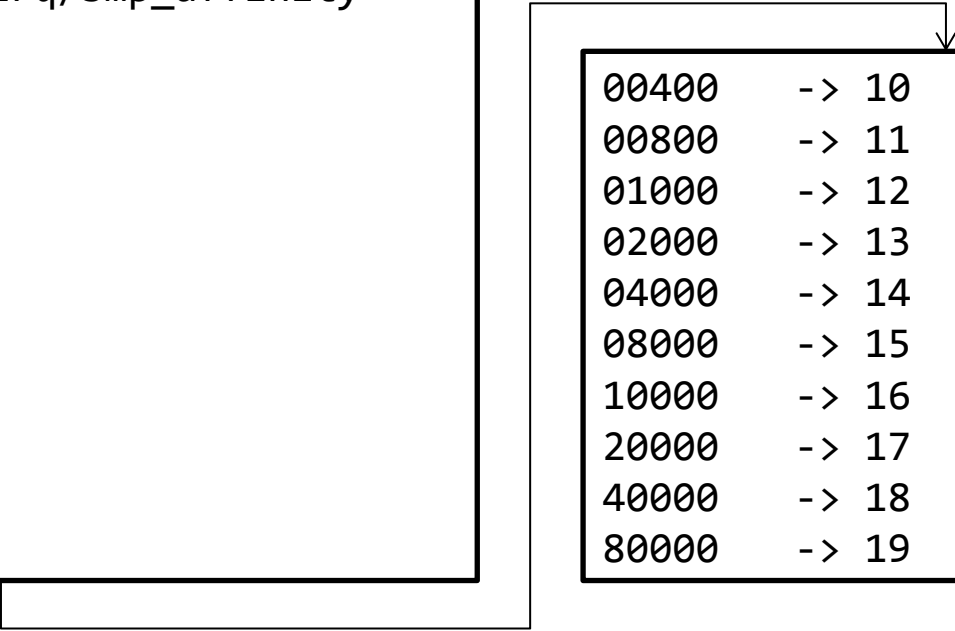
```
00400    -> 10  
00800    -> 11  
01000    -> 12  
02000    -> 13  
04000    -> 14  
08000    -> 15  
10000    -> 16  
20000    -> 17  
40000    -> 18  
80000    -> 19
```

- affinity_hint is evenly distributed
- To honor the hint, add "-h exact" option to irqbalance (via /etc/sysconfig/irqbalance, etc.)*1

Change irqbalance option

- Added "-h exact" and restarted irqbalance

```
$ for ((irq=105; irq<=124; irq++)); do  
>   cat /proc/irq/$irq/smp_affinity  
> done  
00001    -> 0  
00002    -> 1  
00004    -> 2  
00008    -> 3  
00010    -> 4  
00020    -> 5  
00040    -> 6  
00080    -> 7  
00100    -> 8  
00200    -> 9
```



```
00400    -> 10  
00800    -> 11  
01000    -> 12  
02000    -> 13  
04000    -> 14  
08000    -> 15  
10000    -> 16  
20000    -> 17  
40000    -> 18  
80000    -> 19
```

- With hint honored, irqs are distributed to all cores

Change irqbalance option

• sar -u ALL -P ALL 1

20:06:07	CPU	%usr	%nice	%sys	%iowait	%steal	%irq	%soft	%guest	%gnice	%idle
20:06:07	all	0.00	0.00	19.18	0.00	0.00	0.00	80.82	0.00	0.00	0.00
20:06:07	0	0.00	0.00	0.00	0.00	0.00	0.00	100.00	0.00	0.00	0.00
20:06:07	1	0.00	0.00	0.00	0.00	0.00	0.00	100.00	0.00	0.00	0.00
20:06:07	2	0.00	0.00	0.00	0.00	0.00	0.00	100.00	0.00	0.00	0.00
20:06:07	3	0.00	0.00	0.00	0.00	0.00	0.00	100.00	0.00	0.00	0.00
20:06:07	4	0.00	0.00	0.00	0.00	0.00	0.00	100.00	0.00	0.00	0.00
20:06:07	5	0.00	0.00	0.00	0.00	0.00	0.00	100.00	0.00	0.00	0.00
20:06:07	6	0.00	0.00	0.00	0.00	0.00	0.00	100.00	0.00	0.00	0.00
20:06:07	7	0.00	0.00	0.00	0.00	0.00	0.00	100.00	0.00	0.00	0.00
20:06:07	8	0.00	0.00	0.00	0.00	0.00	0.00	100.00	0.00	0.00	0.00
20:06:07	9	0.00	0.00	0.00	0.00	0.00	0.00	100.00	0.00	0.00	0.00
20:06:07	10	0.00	0.00	0.00	0.00	0.00	0.00	100.00	0.00	0.00	0.00
20:06:07	11	0.00	0.00	0.00	0.00	0.00	0.00	100.00	0.00	0.00	0.00
20:06:07	12	0.00	0.00	0.00	0.00	0.00	0.00	100.00	0.00	0.00	0.00
20:06:07	13	0.00	0.00	0.00	0.00	0.00	0.00	100.00	0.00	0.00	0.00
20:06:07	14	0.00	0.00	0.00	0.00	0.00	0.00	100.00	0.00	0.00	0.00
20:06:07	15	0.00	0.00	0.00	0.00	0.00	0.00	100.00	0.00	0.00	0.00
20:06:07	16	0.00	0.00	100.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00
20:06:07	17	0.00	0.00	100.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00
20:06:07	18	0.00	0.00	100.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00
20:06:07	19	0.00	0.00	93.33	0.00	0.00	0.00	6.67	0.00	0.00	0.00

- Though irqstats looks distributed evenly, core 16-19 are not used for softirq...
- Nodes look irrelevant this time

Check rx-queue stats

- **ethtool -S*1**

```
$ ethtool -S ens1f0 | grep 'rx_queue_.*_packets'
rx_queue_0_packets: 198005155
rx_queue_1_packets: 153339750
rx_queue_2_packets: 162870095
rx_queue_3_packets: 172303801
rx_queue_4_packets: 153728776
rx_queue_5_packets: 158138563
rx_queue_6_packets: 164411653
rx_queue_7_packets: 165924489
rx_queue_8_packets: 176545406
rx_queue_9_packets: 165340188
rx_queue_10_packets: 150279834
rx_queue_11_packets: 150983782
rx_queue_12_packets: 157623687
rx_queue_13_packets: 150743910
rx_queue_14_packets: 158634344
rx_queue_15_packets: 158497890
rx_queue_16_packets: 4
rx_queue_17_packets: 3
rx_queue_18_packets: 0
rx_queue_19_packets: 8
```

- **Revealed RSS has not distributed packets to queues 16-19**

RSS Indirection Table

- RSS has indirection table which determines to which queue it spreads packets
- Can be shown by `ethtool -x`

flow hash (hash value from packet header) →

```
$ ethtool -x ens1f0
RX flow hash indirection table for ens1f0 with 20 RX ring(s):
```

0:	0	1	2	3	4	5	6	7
8:	8	9	10	11	12	13	14	15
16:	0	1	2	3	4	5	6	7
24:	8	9	10	11	12	13	14	15
32:	0	1	2	3	4	5	6	7
40:	8	9	10	11	12	13	14	15
48:	0	1	2	3	4	5	6	7
56:	8	9	10	11	12	13	14	15
64:	0	1	2	3	4	5	6	7
72:	8	9	10	11	12	13	14	15
80:	0	1	2	3	4	5	6	7
88:	8	9	10	11	12	13	14	15
96:	0	1	2	3	4	5	6	7
104:	8	9	10	11	12	13	14	15
112:	0	1	2	3	4	5	6	7
120:	8	9	10	11	12	13	14	15

← rx-queue number

- Only rx-queue 0-15 are used, 16-19 not used

RSS Indirection Table

- **Change to use all 0-19?**

```
# ethtool -X ens1f0 equal 20
```

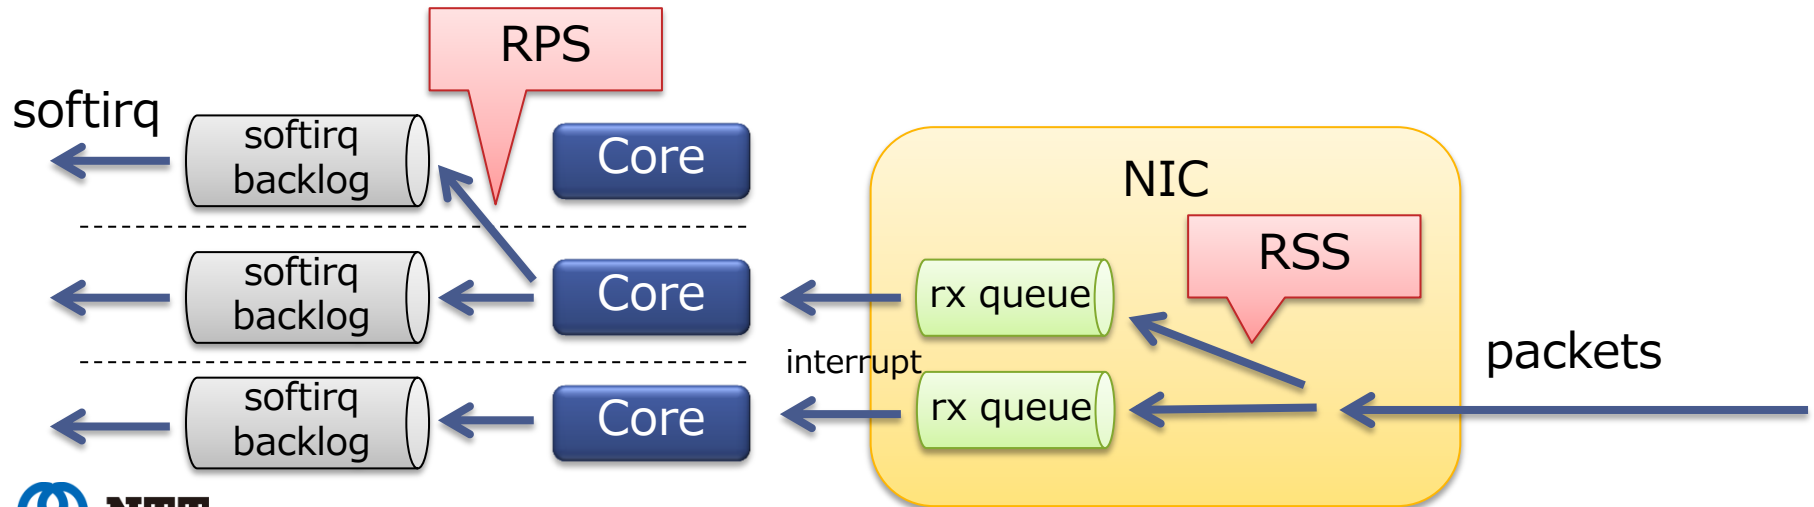
Cannot set RX flow hash configuration: Invalid argument

- **This NIC's max rx-queues in the indirection table is actually 16 so we cannot use 20 queues**

- although we have 64 rx-queues...

- **Use RPS instead**

- Software emulation of RSS



- **This time I spread flows from rx-queue 6-9 to core 6-9 and 16-19**
 - Because they are all in Node 1
 - rx-queue 6 -> core 6, 16
 - rx-queue 7 -> core 7, 17
 - rx-queue 8 -> core 8, 18
 - rx-queue 9 -> core 9, 19

```
# echo 10040 > /sys/class/net/ens1f0/queues/rx-6/rps_cpus  
# echo 20080 > /sys/class/net/ens1f0/queues/rx-7/rps_cpus  
# echo 40100 > /sys/class/net/ens1f0/queues/rx-8/rps_cpus  
# echo 80200 > /sys/class/net/ens1f0/queues/rx-9/rps_cpus
```

Use RPS

- **sar -u ALL -P ALL 1**

20:18:53	CPU	%usr	%nice	%sys	%iowait	%steal	%irq	%soft	%guest	%gnice	%idle
20:18:54	all	0.00	0.00	2.38	0.00	0.00	0.00	97.62	0.00	0.00	0.00
20:18:54	0	0.00	0.00	0.00	0.00	0.00	0.00	100.00	0.00	0.00	0.00
20:18:54	1	0.00	0.00	0.00	0.00	0.00	0.00	100.00	0.00	0.00	0.00
20:18:54	2	0.00	0.00	0.00	0.00	0.00	0.00	100.00	0.00	0.00	0.00
20:18:54	3	0.00	0.00	0.00	0.00	0.00	0.00	100.00	0.00	0.00	0.00
20:18:54	4	0.00	0.00	0.00	0.00	0.00	0.00	100.00	0.00	0.00	0.00
20:18:54	5	0.00	0.00	0.00	0.00	0.00	0.00	100.00	0.00	0.00	0.00
20:18:54	6	0.00	0.00	0.00	0.00	0.00	0.00	100.00	0.00	0.00	0.00
20:18:54	7	0.00	0.00	0.00	0.00	0.00	0.00	100.00	0.00	0.00	0.00
20:18:54	8	0.00	0.00	0.00	0.00	0.00	0.00	100.00	0.00	0.00	0.00
20:18:54	9	0.00	0.00	0.00	0.00	0.00	0.00	100.00	0.00	0.00	0.00
20:18:54	10	0.00	0.00	0.00	0.00	0.00	0.00	100.00	0.00	0.00	0.00
20:18:54	11	0.00	0.00	0.00	0.00	0.00	0.00	100.00	0.00	0.00	0.00
20:18:54	12	0.00	0.00	0.00	0.00	0.00	0.00	100.00	0.00	0.00	0.00
20:18:54	13	0.00	0.00	0.00	0.00	0.00	0.00	100.00	0.00	0.00	0.00
20:18:54	14	0.00	0.00	0.00	0.00	0.00	0.00	100.00	0.00	0.00	0.00
20:18:54	15	0.00	0.00	0.00	0.00	0.00	0.00	100.00	0.00	0.00	0.00
20:18:54	16	0.00	0.00	15.56	0.00	0.00	0.00	84.44	0.00	0.00	0.00
20:18:54	17	0.00	0.00	6.98	0.00	0.00	0.00	93.02	0.00	0.00	0.00
20:18:54	18	0.00	0.00	18.18	0.00	0.00	0.00	81.82	0.00	0.00	0.00
20:18:54	19	2.27	0.00	6.82	0.00	0.00	0.00	90.91	0.00	0.00	0.00

- **softirq is almost evenly distributed**

RSS & affinity_hint & RPS

- Now thanks to affinity_hint and RPS, we succeeded to spread flows almost evenly
- Performance change
 - Before: **270,000** tps (approx. **360Mbps**)
 - After: **17,000** tps (approx. **23Mbps**)
 - Got worse...
- Probably the reason is too heavy softirq
 - softirq is almost 100% in total
 - Need finer-grained profiling than sar

- **perf**

- Profiling tool developed in kernel tree
- Identify hot spots by sampling CPU cycles

- **Example usage of perf**

- `perf record -a -g -- sleep 5`
 - Save sampling results for 5 seconds to perf.data file

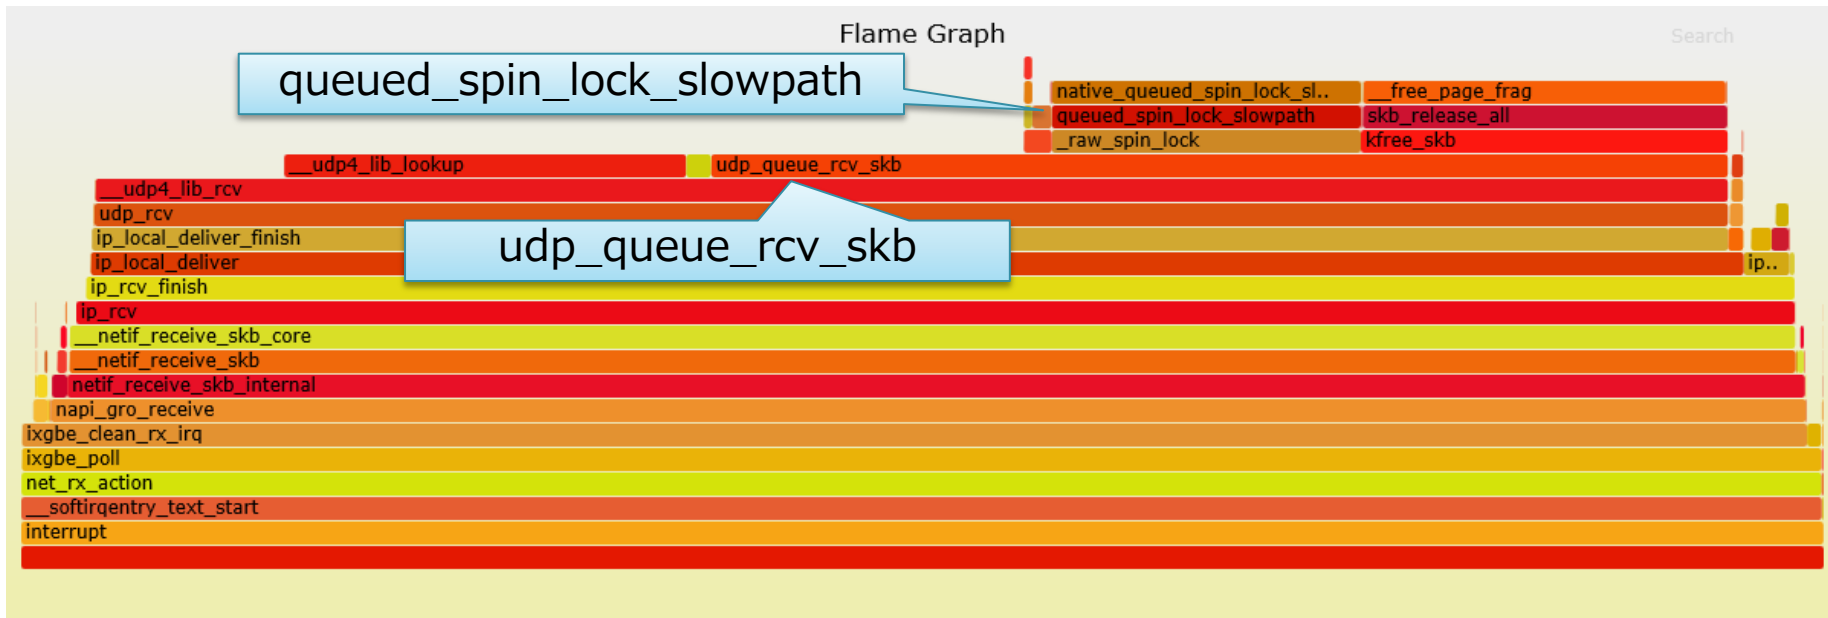
- **FlameGraph**

- Visualize perf.data in svg format
- <https://github.com/brendangregg/FlameGraph>

Profile softirq

- **FlameGraph of CPU0*1**

- x-axis (width): CPU consumption
- y-axis (height): Depth of call stack

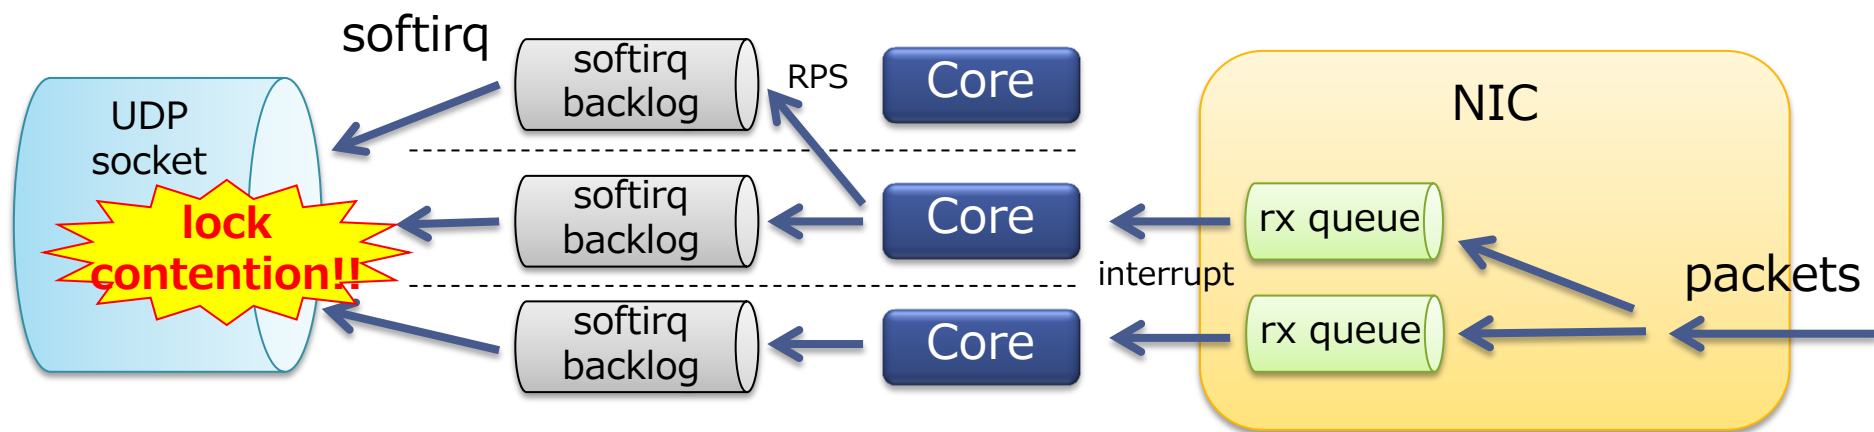


- **queued_spin_lock_slowpath: lock is contended**
- **udp_queue_rcv_skb: acquires socket lock**

*1 I filtered call stack under irq context from output of perf script to make the chart easier to see
irq context is shown as "interrupt" here

Socket lock contention

- Echo server has only one socket bound to a certain port
- softirq of each core pushes packets into socket queue concurrently

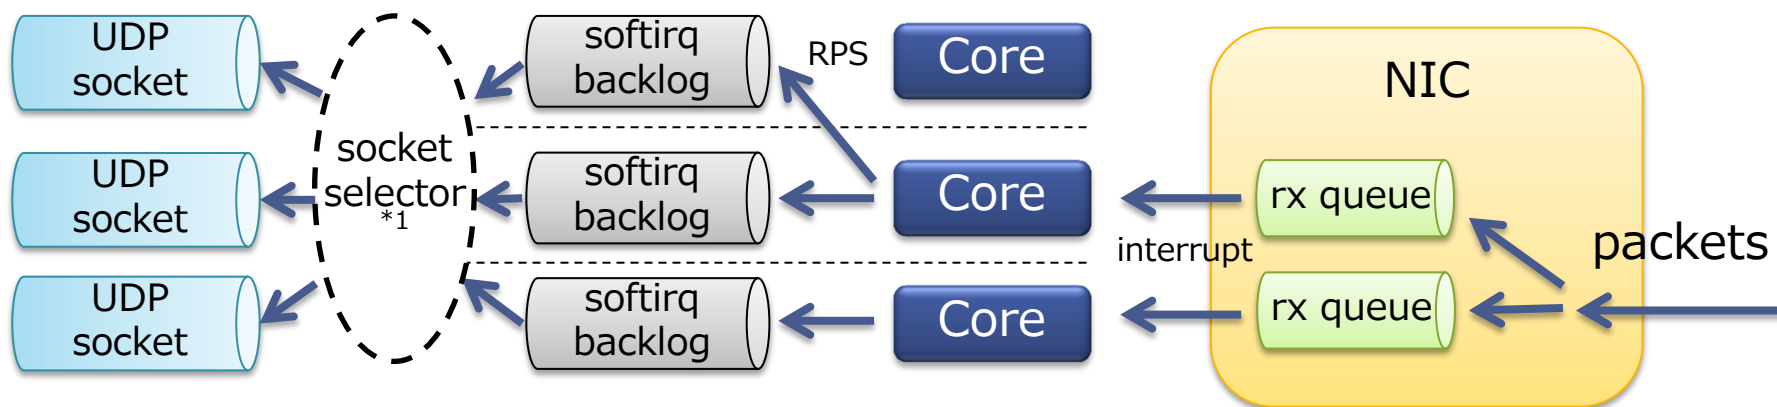


- socket lock gets contended

Avoid lock contention

• Split sockets by SO_REUSEPORT

- Introduced by kernel 3.9



• SO_REUSEPORT allows multiple UDP sockets to bind the same port

- One of the sockets is chosen on queueing each packet

```
int on = 1;
int sock = socket(AF_INET, SOCK_DGRAM, 0);
setsockopt(sock, SOL_SOCKET, SO_REUSEPORT, &on, sizeof(on));
bind(sock, ...);
```

*1 select a socket by flow (packet header) hash by default

Use SO_REUSEPORT

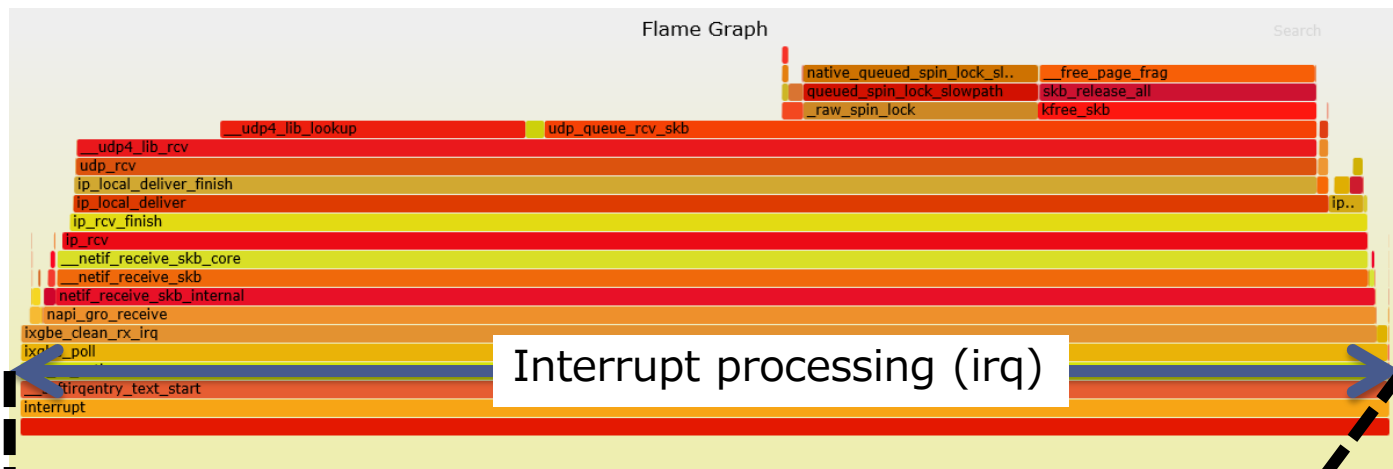
- **sar -u ALL -P ALL 1**

20:44:33	CPU	%usr	%nice	%sys	%iowait	%steal	%irq	%soft	%guest	%gnice	%idle
20:44:34	all	3.26	0.00	37.23	0.00	0.00	0.00	59.52	0.00	0.00	0.00
20:44:34	0	3.33	0.00	28.33	0.00	0.00	0.00	68.33	0.00	0.00	0.00
20:44:34	1	3.33	0.00	25.00	0.00	0.00	0.00	71.67	0.00	0.00	0.00
20:44:34	2	1.67	0.00	23.33	0.00	0.00	0.00	75.00	0.00	0.00	0.00
20:44:34	3	3.28	0.00	32.79	0.00	0.00	0.00	63.93	0.00	0.00	0.00
20:44:34	4	3.33	0.00	33.33	0.00	0.00	0.00	63.33	0.00	0.00	0.00
20:44:34	5	1.69	0.00	23.73	0.00	0.00	0.00	74.58	0.00	0.00	0.00
20:44:34	6	3.28	0.00	50.82	0.00	0.00	0.00	45.90	0.00	0.00	0.00
20:44:34	7	3.45	0.00	50.00	0.00	0.00	0.00	46.55	0.00	0.00	0.00
20:44:34	8	1.69	0.00	37.29	0.00	0.00	0.00	61.02	0.00	0.00	0.00
20:44:34	9	1.67	0.00	33.33	0.00	0.00	0.00	65.00	0.00	0.00	0.00
20:44:34	10	1.69	0.00	18.64	0.00	0.00	0.00	79.66	0.00	0.00	0.00
20:44:34	11	3.23	0.00	35.48	0.00	0.00	0.00	61.29	0.00	0.00	0.00
20:44:34	12	1.69	0.00	27.12	0.00	0.00	0.00	71.19	0.00	0.00	0.00
20:44:34	13	1.67	0.00	21.67	0.00	0.00	0.00	76.67	0.00	0.00	0.00
20:44:34	14	1.67	0.00	21.67	0.00	0.00	0.00	76.67	0.00	0.00	0.00
20:44:34	15	3.33	0.00	35.00	0.00	0.00	0.00	61.67	0.00	0.00	0.00
20:44:34	16	6.67	0.00	68.33	0.00	0.00	0.00	25.00	0.00	0.00	0.00
20:44:34	17	5.00	0.00	65.00	0.00	0.00	0.00	30.00	0.00	0.00	0.00
20:44:34	18	6.78	0.00	54.24	0.00	0.00	0.00	38.98	0.00	0.00	0.00
20:44:34	19	4.92	0.00	63.93	0.00	0.00	0.00	31.15	0.00	0.00	0.00


- **CPU consumption in softirq became some more reasonable**

Use SO_REUSEPORT

• before



• after



Userspace starts to work

Use SO_REUSEPORT




- **Performance change**

- RSS: 270,000 tps (approx. 360Mbps)
- +affinity_hint+RPS: 17,000 tps (approx. 23Mbps)
- +SO_REUSEPORT: **2,540,000** tps (approx. **3370Mbps**)
 - Great improvement!
 - but...

Use SO_REUSEPORT

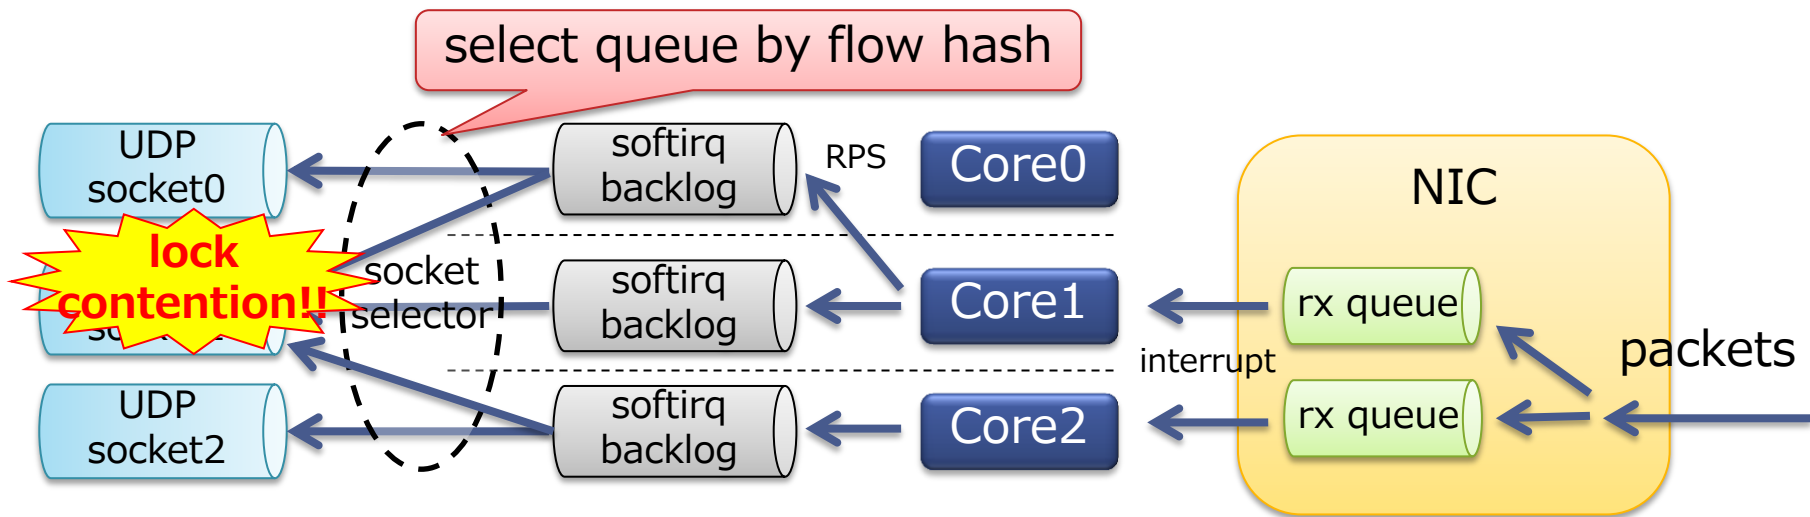
- More analysis



- Socket lock is still contended

Socket lock contention again

- **SO_REUSEPORT** uses flow hash to select queue by default
- Same sockets can be selected by different cores

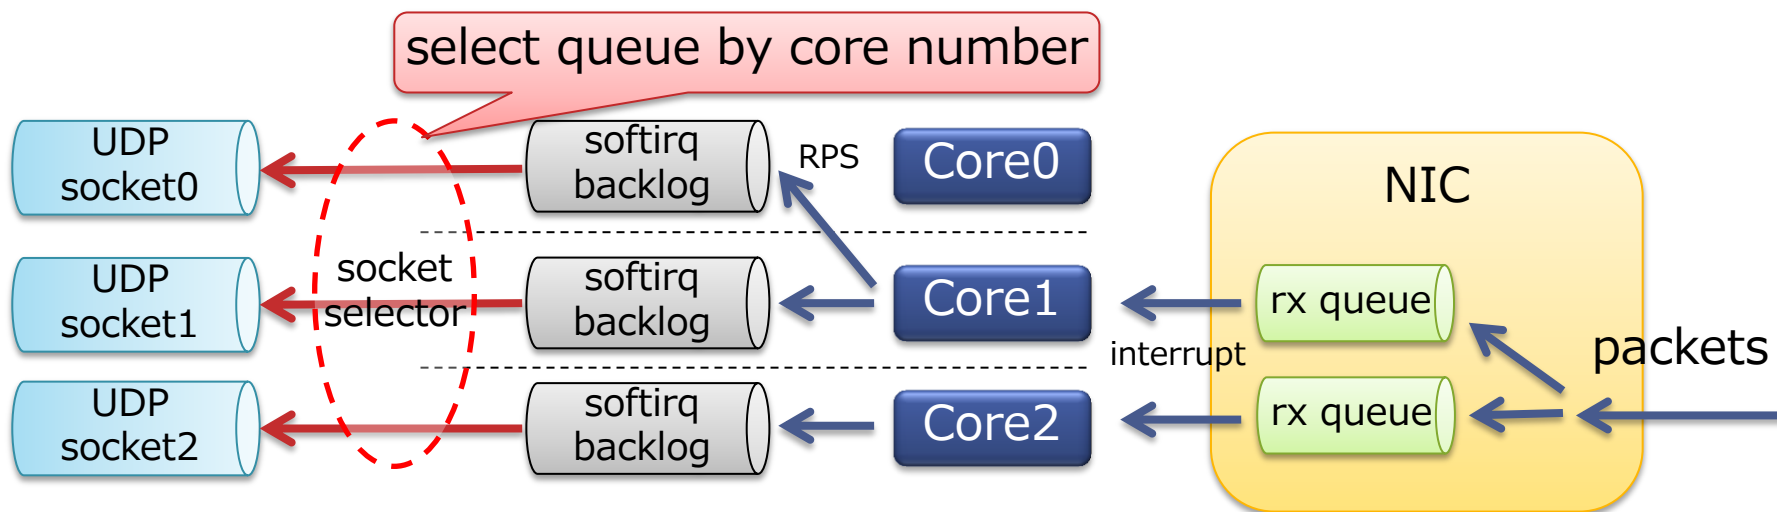


- **Socket lock still gets contended**

Avoid socket lock contention

• Select socket by core number

- Realized by `SO_ATTACH_REUSEPORT_CBPF/EBPF`^{*1}
- Introduced by kernel 4.5




• No lock contention between softirq

• Usage

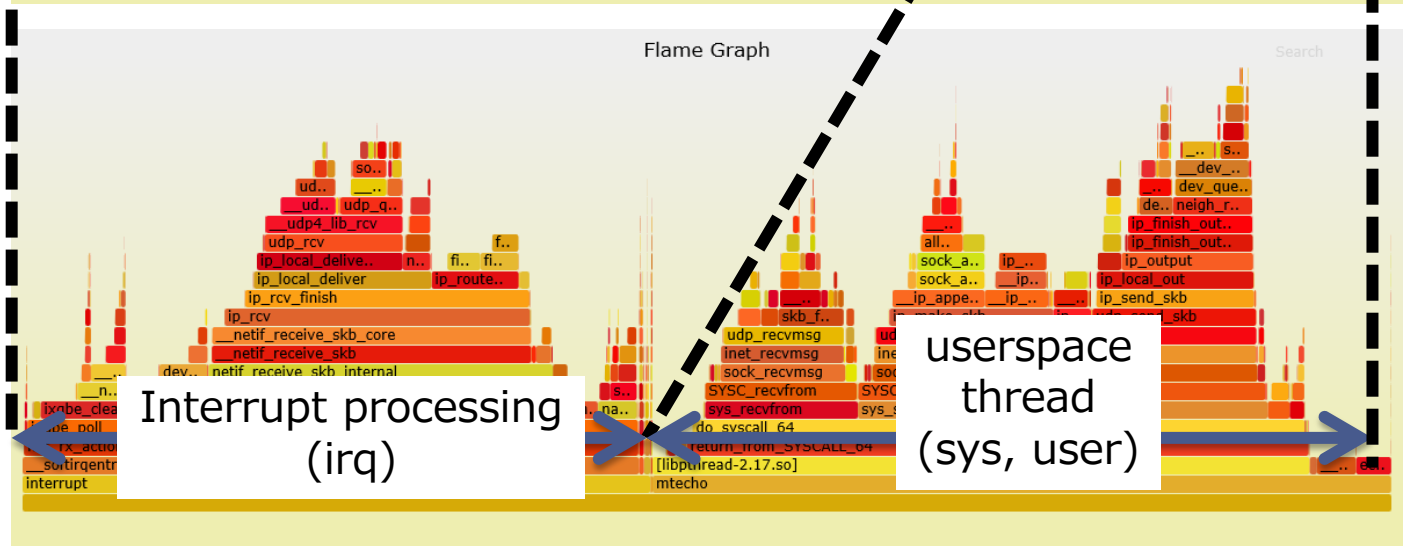
- See example in kernel source tree
 - `tools/testing/selftests/net/reuseport_bpf_cpu.c`

Use SO_ATTACH_REUSEPORT_EPBF

- before



- after



irq overhead gets less

Use SO_ATTACH_REUSEPORT_EBPF

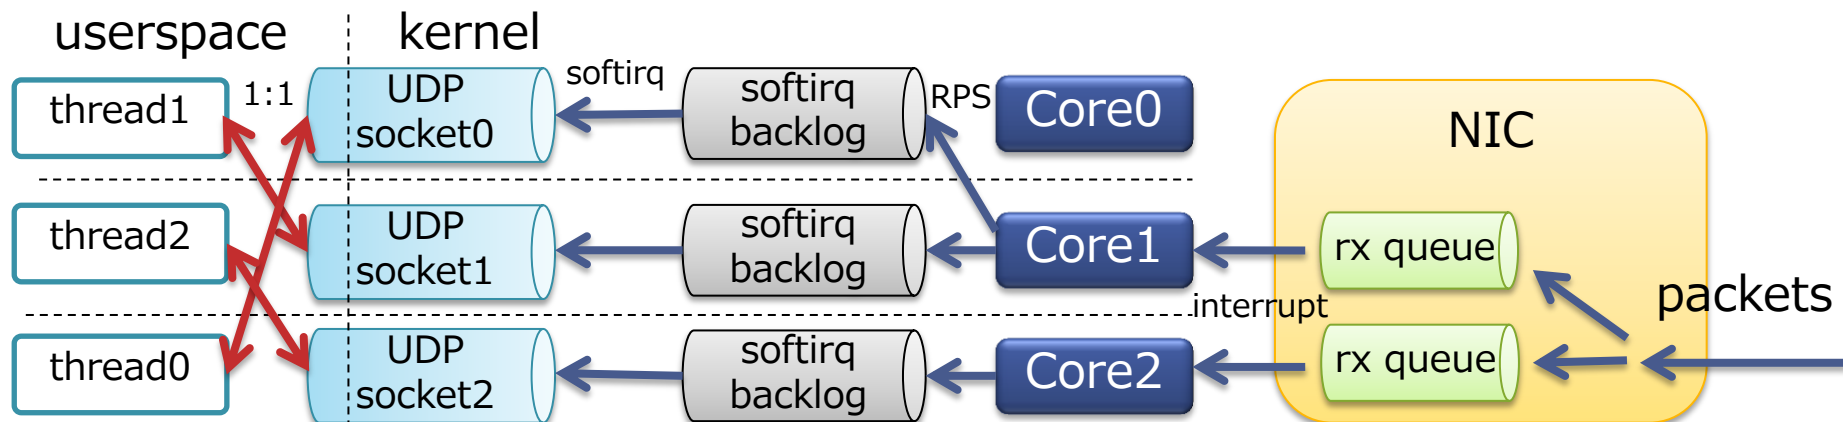


- **Performance change**

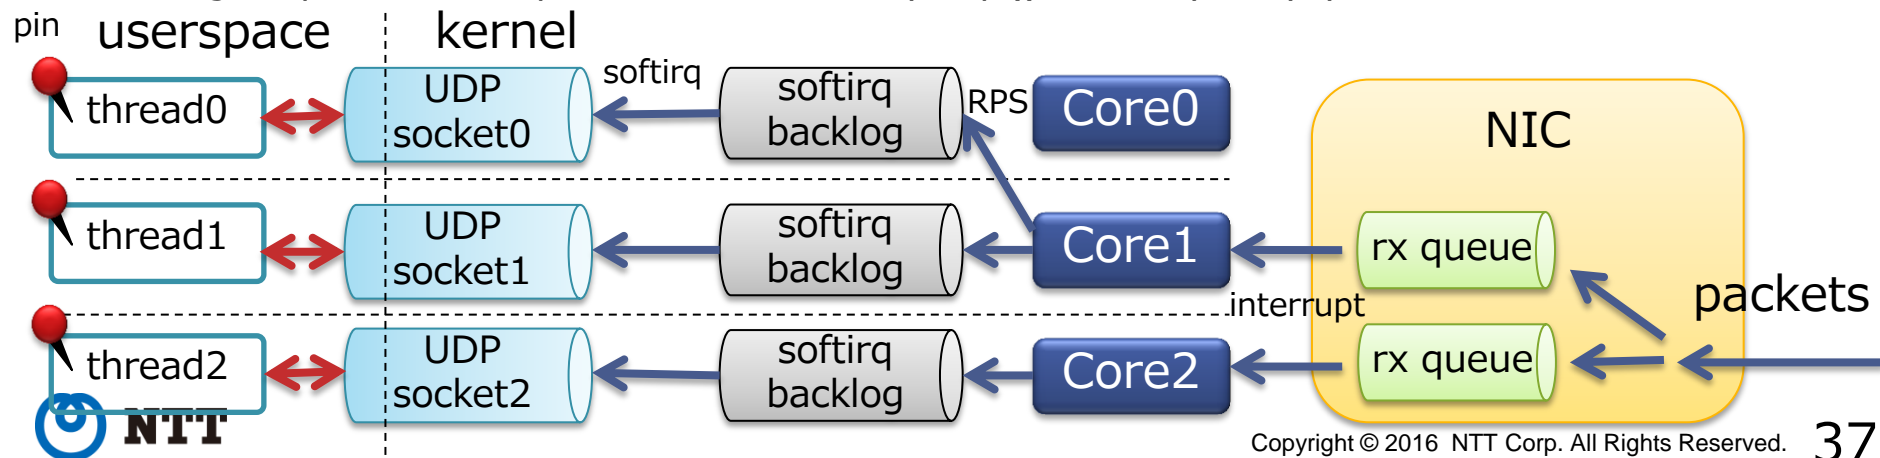
- RSS: 270,000 tps (approx. 360Mbps)
- +affinity_hint+RPS: 17,000 tps (approx. 23Mbps)
- +SO_REUSEPORT: 2,540,000 tps (approx. 3370Mbps)
- +SO_ATTACH_...: **4,250,000 tps (approx. 5640Mbps)**

Pin userspace threads

- **Userspace threads : sockets == 1 : 1**
 - No lock contention
- **But not necessarily on the same core as softirq**



- **Pin userspace thread on the same core for better cache affinity**
 - cgroup, taskset, pthread_setaffinity_np(), ... any way you like



Pin userspace threads



• Performance change

- RSS: 270,000 tps (approx. 360Mbps)
- +affinity_hint+RPS: 17,000 tps (approx. 23Mbps)
- +SO_REUSEPORT: 2,540,000 tps (approx. 3370Mbps)
- +SO_ATTACH_...: 4,250,000 tps (approx. 5640Mbps)
- +Pin threads: **5,050,000 tps (approx. 6710Mbps)**

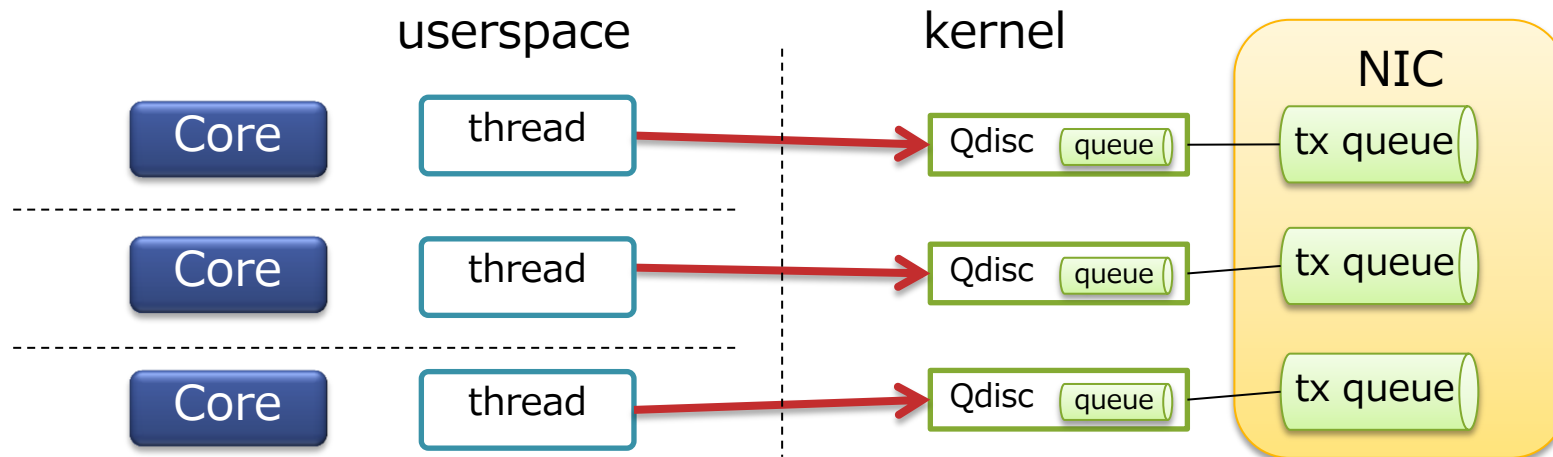
Tx lock contention?



- So far everything has been about Rx
- No lock contention on Tx?

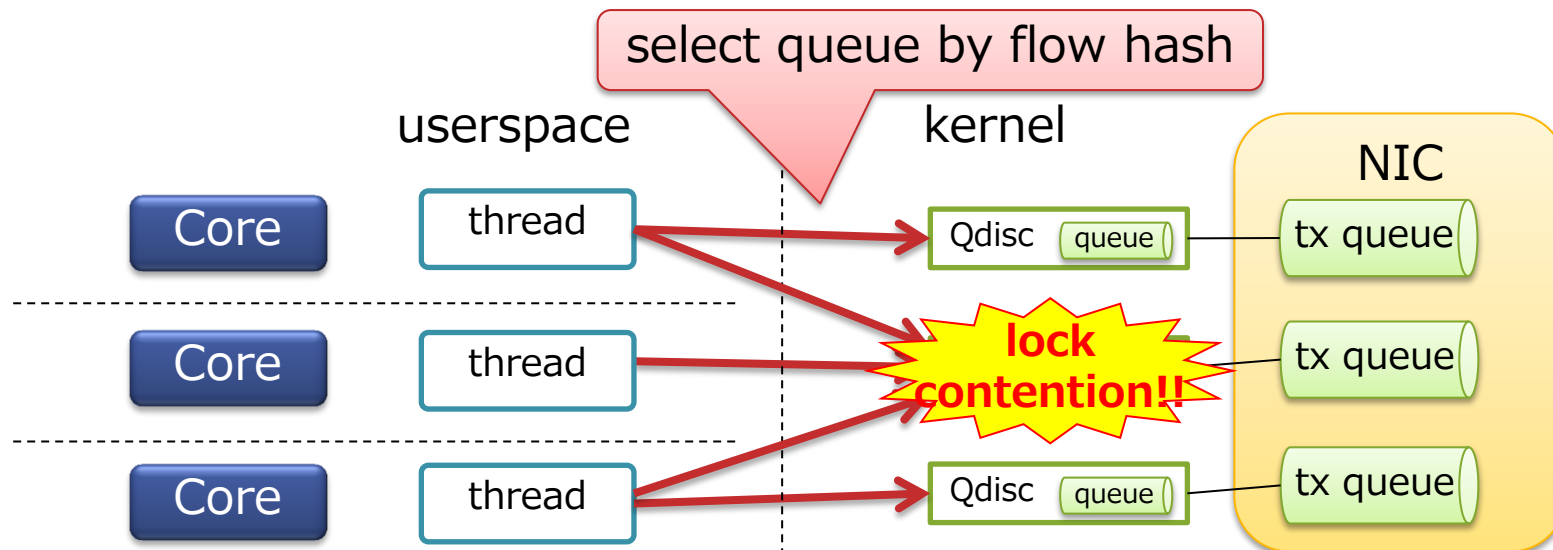
Tx queue

- kernel has Qdisc (Queueing discipline)
- Each Qdisc is linked to NIC tx-queue
- Each Qdisc has its lock



Tx queue lock contention

- By default Qdisc is selected by flow hash
- Thus lock contention can happen

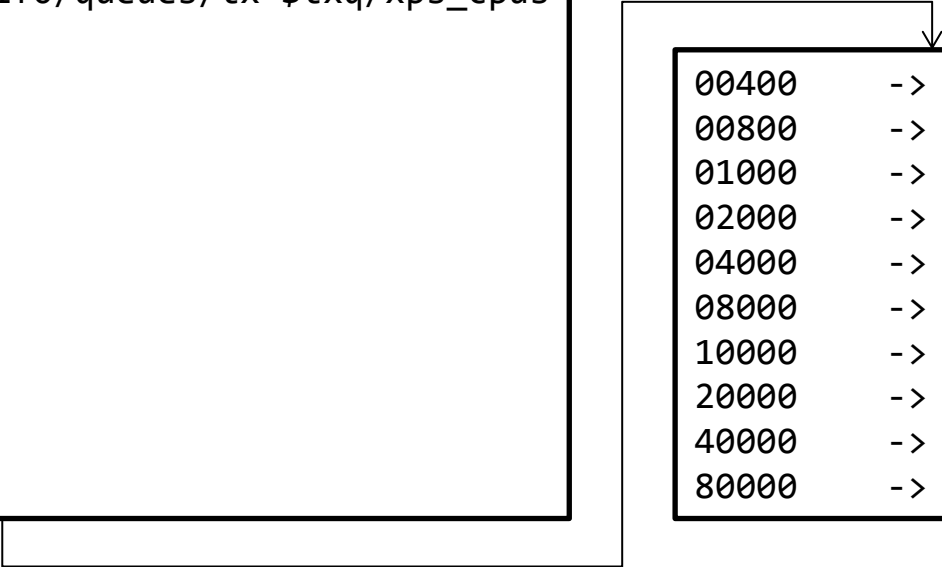


- We haven't seen contention on Tx, why?

Avoid Tx queue lock contention

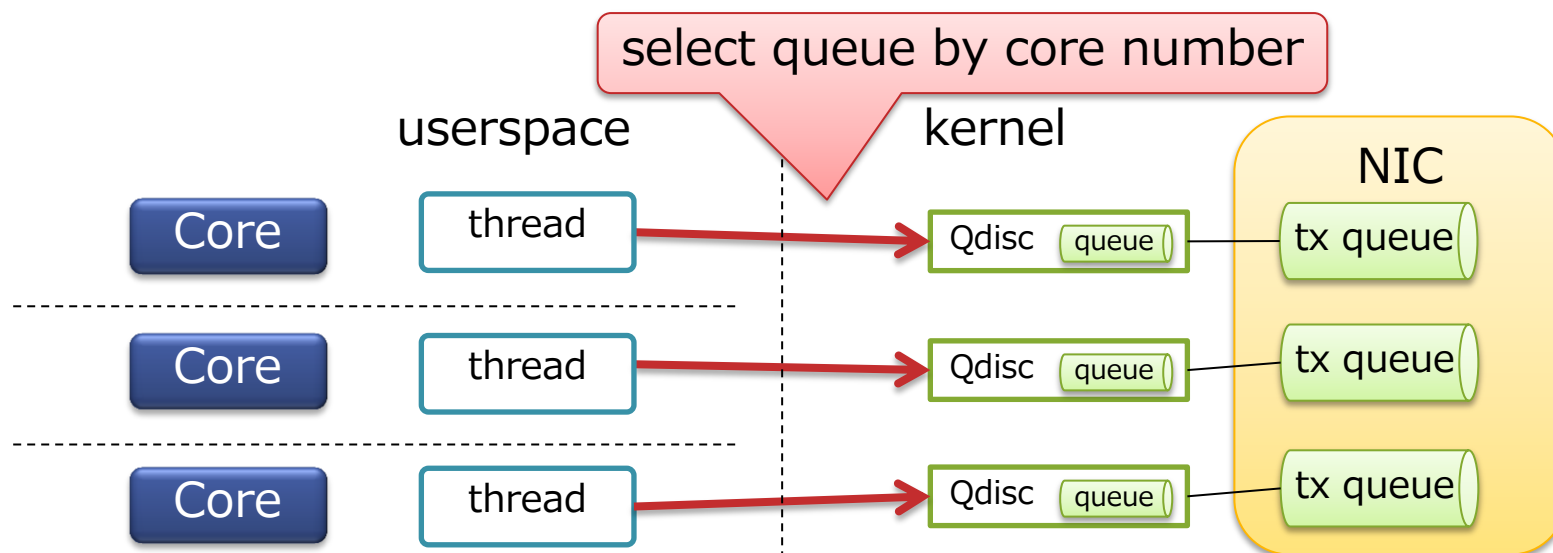
- Because ixgbe (Intel 10GbE NIC driver) has an ability to set XPS automatically

```
$ for ((txq=0; txq<20; txq++)); do  
>   cat /sys/class/net/ens1f0/queues/tx-$txq/xps_cpus  
> done  
00001    -> core 0  
00002    -> core 1  
00004    -> core 2  
00008    -> core 3  
00010    -> core 4  
00020    -> core 5  
00040    -> core 6  
00080    -> core 7  
00100    -> core 8  
00200    -> core 9
```



```
00400    -> core 10  
00800    -> core 11  
01000    -> core 12  
02000    -> core 13  
04000    -> core 14  
08000    -> core 15  
10000    -> core 16  
20000    -> core 17  
40000    -> core 18  
80000    -> core 19
```

- XPS allows kernel to select Tx queue (Qdisc) by core number



- Tx has no lock contention

How effective is XPS?

- Try disabling it

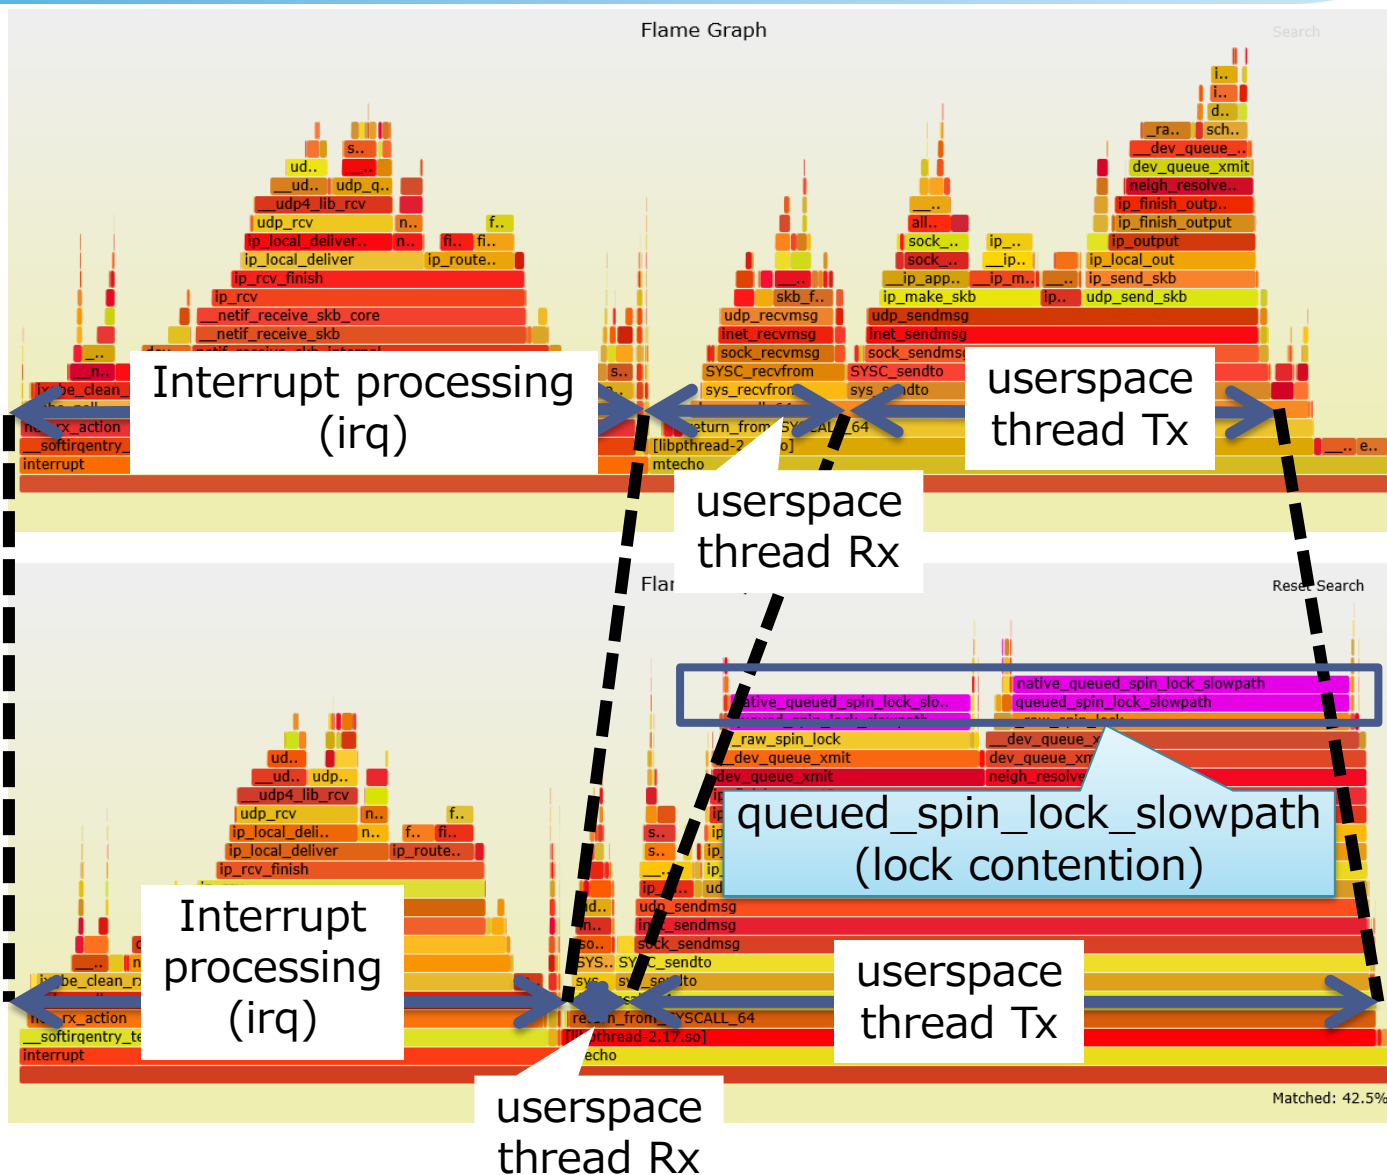
```
# for ((txq=0; txq<20; txq++)); do  
>   echo 0 > /sys/class/net/ens1f0/queues/tx-$txq/xps_cpus  
> done
```

- Before: **5,050,000** tps (approx. **6710Mbps**)
- After: **1,086,000** tps (approx. **1440Mbps**)

Disabling XPS

- XPS enabled

- XPS disabled



Enable XPS

- **Enable XPS again**

```
# echo 00001 > /sys/class/net/<NIC>/queues/tx-0/xps_cpus  
# echo 00002 > /sys/class/net/<NIC>/queues/tx-1/xps_cpus  
# echo 00004 > /sys/class/net/<NIC>/queues/tx-2/xps_cpus  
# echo 00008 > /sys/class/net/<NIC>/queues/tx-3/xps_cpus  
...
```


- **Although ixgbe can automatically set XPS, not all drivers can do that**
- **Make sure to check xps_cpus is configured**

Optimization per core



- By making full use of multi-core with avoiding contention, we achieved
 - **5,050,000** tps (approx. **6710Mbps**)
- To get more performance, reduce overhead per core

Optimization per core



- GRO is enabled by default
- Consuming 3.57% of CPU time

- **GRO is not applicable to UDP^{*1}**
- **Disable it for UDP servers**

```
# ethtool -K <NIC> gro off
```

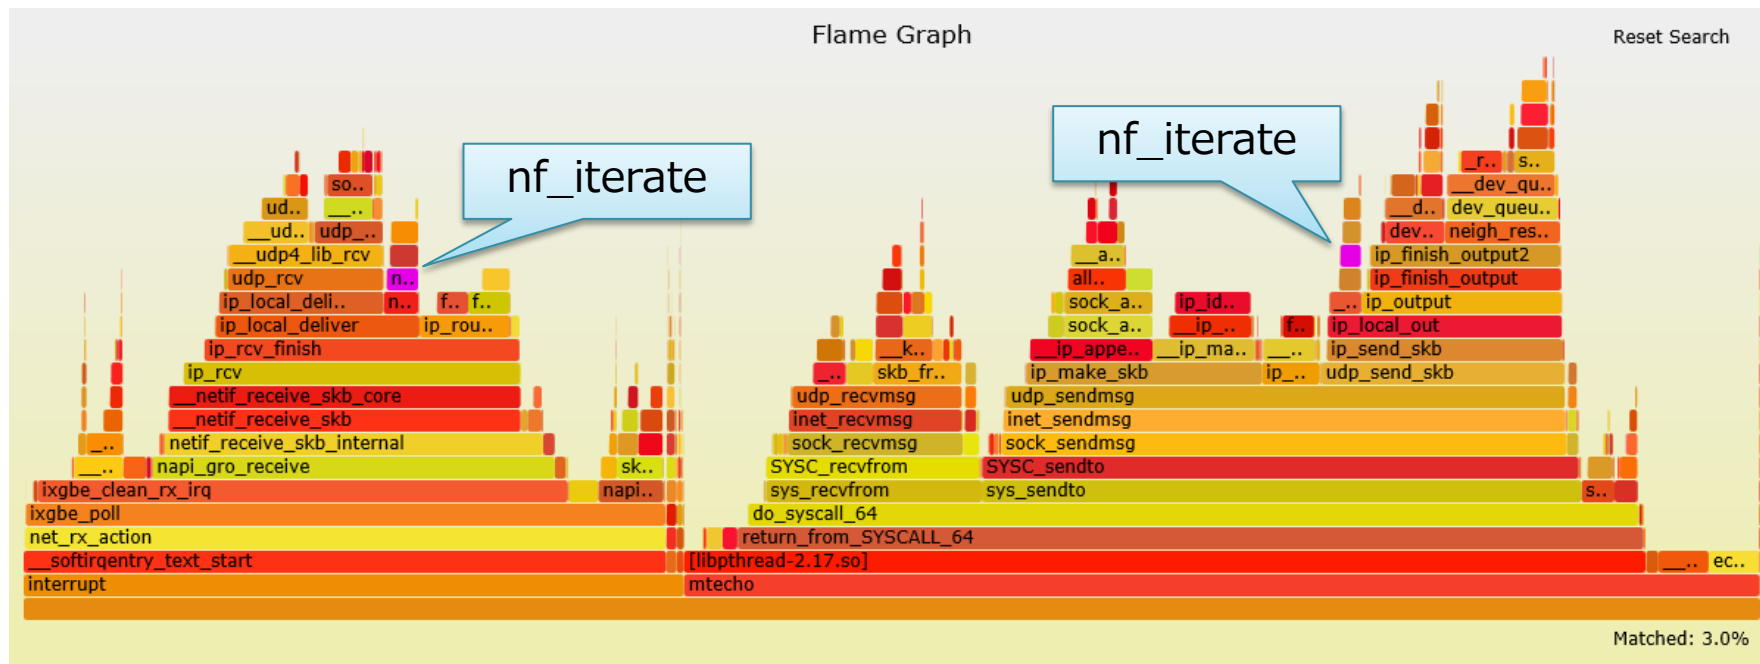
- **WARNING:**
 - Don't disable it if TCP performance matters
 - Disabling GRO makes TCP rx throughput miserably low
 - Don't disable it on KVM hypervisors as well
 - GRO boost throughput of tunneling protocol traffic as well as guest's TCP traffic on hypervisors

Disable GRO

• Performance change

- RSS (+XPS): 270,000 tps (approx. 360Mbps)
- +affinity_hint+RPS: 17,000 tps (approx. 23Mbps)
- +SO_REUSEPORT: 2,540,000 tps (approx. 3370Mbps)
- +SO_ATTACH_...: 4,250,000 tps (approx. 5640Mbps)
- +Pin threads: 5,050,000 tps (approx. 6710Mbps)
- +Disable GRO: **5,180,000 tps (approx. 6880Mbps)**

Optimization per core



- **iptables-related processing (nf_iterate) is performed**
 - Although I have not added any rule to iptables
- **Consuming 3.00% of CPU time**

iptables (netfilter)

- With iptables kernel module loaded, even if you don't have any rules, it can incur some overhead
- Some distributions load iptables module even when you don't add any rule
- If you are not using iptables, unload the module


```
# modprobe -r iptable_filter  
# modprobe -r ip_tables
```

Unload iptables



• Performance change

- RSS (+XPS): 270,000 tps (approx. 360Mbps)
- +affinity_hint+RPS: 17,000 tps (approx. 23Mbps)
- +SO_REUSEPORT: 2,540,000 tps (approx. 3370Mbps)
- +SO_ATTACH_...: 4,250,000 tps (approx. 5640Mbps)
- +Pin threads: 5,050,000 tps (approx. 6710Mbps)
- +Disable GRO: 5,180,000 tps (approx. 6880Mbps)
- +Unload iptables: **5,380,000 tps (approx. 7140Mbps)**



FIB lookup on Rx

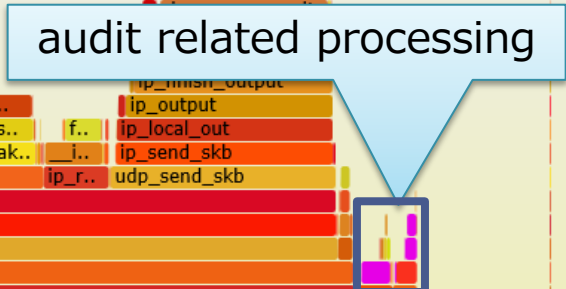
- **One of two times of table lookup is for validating source IP addresses**
 - Reverse path filter
 - Local address check
- **If you really don't need source validation, you can skip it**

```
# sysctl -w net.ipv4.conf.all.rp_filter=0  
# sysctl -w net.ipv4.conf.<NIC>.rp_filter=0  
# sysctl -w net.ipv4.conf.all.accept_local=1
```

Disable source validation

• Performance change

- RSS (+XPS): 270,000 tps (approx. 360Mbps)
- +affinity_hint+RPS: 17,000 tps (approx. 23Mbps)
- +SO_REUSEPORT: 2,540,000 tps (approx. 3370Mbps)
- +SO_ATTACH_...: 4,250,000 tps (approx. 5640Mbps)
- +Pin threads: 5,050,000 tps (approx. 6710Mbps)
- +Disable GRO: 5,180,000 tps (approx. 6880Mbps)
- +Unload iptables: 5,380,000 tps (approx. 7140Mbps)
- +Disable validation: **5,490,000** tps (approx. **7290Mbps**)



- 


- **If you don't need audit, disable it**

```
# systemctl disable auditd  
# reboot
```

• Performance change

- RSS (+XPS): 270,000 tps (approx. 360Mbps)
- +affinity_hint+RPS: 17,000 tps (approx. 23Mbps)
- +SO_REUSEPORT: 2,540,000 tps (approx. 3370Mbps)
- +SO_ATTACH_...: 4,250,000 tps (approx. 5640Mbps)
- +Pin threads: 5,050,000 tps (approx. 6710Mbps)
- +Disable GRO: 5,180,000 tps (approx. 6880Mbps)
- +Unload iptables: 5,380,000 tps (approx. 7140Mbps)
- +Disable validation: 5,490,000 tps (approx. 7290Mbps)
- +Disable audit: **5,860,000 tps (approx. 7780Mbps)**

Optimization per core



- IP ID field calculation (__ip_select_ident) is heavy
- Consuming 4.82% of CPU time

IP ID field calculation



- **This is an environment-specific issue**

- This happens if many clients has the same IP address
 - Cache contention by atomic operations
- It is very likely you don't see this amount of CPU consumption without using tunneling protocol

- **If you really see this problem...**

- You can skip it only if you never send over-mtu-sized packets
 - Though it is very strict

```
int pmtu = IP_PMTUDISC_D0;  
setsockopt(sock, IPPROTO_IP, IP_MTU_DISCOVER, &pmtu, sizeof(pmtu));
```

Skip IP ID calculation

• Performance change

- RSS (+XPS): 270,000 tps (approx. 360Mbps)
- +affinity_hint+RPS: 17,000 tps (approx. 23Mbps)
- +SO_REUSEPORT: 2,540,000 tps (approx. 3370Mbps)
- +SO_ATTACH_...: 4,250,000 tps (approx. 5640Mbps)
- +Pin threads: 5,050,000 tps (approx. 6710Mbps)
- +Disable GRO: 5,180,000 tps (approx. 6880Mbps)
- +Unload iptables: 5,380,000 tps (approx. 7140Mbps)
- +Disable validation: 5,490,000 tps (approx. 7290Mbps)
- +Disable audit: 5,860,000 tps (approx. 7780Mbps)
- +Skip ID calculation: **6,010,000** tps (approx. **7980Mbps**)

Hyper threading

- **So far we have not enabled hyper threading**
- **It makes the number of logical cores 40**
 - Number of physical cores are 20 in this box
- **With 40 cores we need to rely on RPS more**
 - Remind: Max usable rx-queues == 16
- **Enable hyper-threading and set RPS on all rx-queues**
 - queue 0 -> core 0, 20
 - queue 1 -> core 1, 21
 - ...
 - queue 10 -> core 10, 16, 30
 - queue 11 -> core 11, 17, 31
 - ...


Hyper threading

• Performance change

• RSS (+XPS):	270,000 tps (approx.	360Mbps)
• +affinity_hint+RPS:	17,000 tps (approx.	23Mbps)
• +SO_REUSEPORT:	2,540,000 tps (approx.	3370Mbps)
• +SO_ATTACH_...:	4,250,000 tps (approx.	5640Mbps)
• +Pin threads:	5,050,000 tps (approx.	6710Mbps)
• +Disable GRO:	5,180,000 tps (approx.	6880Mbps)
• +Unload iptables:	5,380,000 tps (approx.	7140Mbps)
• +Disable validation:	5,490,000 tps (approx.	7290Mbps)
• +Disable audit:	5,860,000 tps (approx.	7780Mbps)
• +Skip ID calculation:	6,010,000 tps (approx.	7980Mbps)
• +Hyper threading:	7,010,000 tps (approx.	9310Mbps)


- I guess more rx-queues would realize even better performance number

More hot spots



- Tx Qdisc lock (`_raw_spin_lock`) is heavy
- Not contended but involves many atomic operations
- Being optimized in Linux netdev community

More hot spots



- Memory alloc/free (slab)
- Being optimized in netdev community as well

- **Virtualization**

- UDP servers as guests
- Hypervisor can saturate CPUs or drop packets
- We are going to investigate ways to boost performance in virtualized environment as well

Summary

- **For 100bytes, we can achieve almost 10G**
 - From: 270,000 tps (approx. 360Mbps)
 - To: 7,010,000 tps (approx. 9310Mbps)
 - Of course we need to take into account additional userspace work in real applications so this number is not applicable as is
- **To boost UDP performance**
 - Applications (Most important!)
 - implement SO_REUSEPORT
 - implement SO_ATTACH_REUSEPORT_EBPF/CBPF
 - These are useful for TCP listening sockets as well
 - OS settings
 - Use RPS if rx-queues are not enough
 - Make sure XPS is configured
 - Consider other tunings to reduce per-core overhead
 - Disable GRO
 - Unload iptables
 - Disable source IP validation
 - Disable auditd
 - Hardware
 - Use NICs which have enough RSS rx-queues if possible (as many queues as core num)

Thank you!